

Spark Day 3 Assignment

1. Pipeline 1 – RDBMS → Spark → Amazon Keyspaces (Cassandra)

Purpose: Migrate and denormalize relational data from RDS MySQL to Amazon Keyspaces for analytics and high-performance queries.

Architecture:

- Source: AWS RDS MySQL (3 normalized tables: customers, orders, order_items)
- Processing: Apache Spark (Scala) with JDBC and Cassandra Connector
- Target: Amazon Keyspaces table sales_data in keyspace spark_keyspace

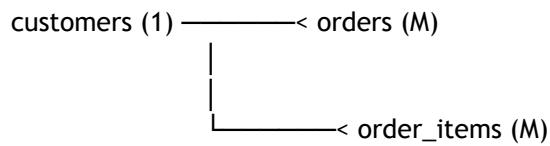
Database Schema:

```
CREATE TABLE customers (
    customer_id INT UNSIGNED AUTO_INCREMENT PRIMARY KEY,
    name        VARCHAR(100) NOT NULL,
    email       VARCHAR(150) UNIQUE NOT NULL,
    city        VARCHAR(80) NOT NULL
);

CREATE TABLE orders (
    order_id    INT UNSIGNED AUTO_INCREMENT PRIMARY KEY,
    customer_id INT UNSIGNED NOT NULL,
    order_date  DATE NOT NULL,
    amount      DECIMAL(12,2) NOT NULL,
    CONSTRAINT fk_orders_customer
        FOREIGN KEY (customer_id) REFERENCES customers(customer_id)
        ON DELETE RESTRICT ON UPDATE CASCADE
);

CREATE TABLE order_items (
    item_id     INT UNSIGNED AUTO_INCREMENT PRIMARY KEY,
    order_id    INT UNSIGNED NOT NULL,
    product_name VARCHAR(200) NOT NULL,
    quantity    INT UNSIGNED NOT NULL DEFAULT 1,
    CONSTRAINT fk_order_items_order
        FOREIGN KEY (order_id) REFERENCES orders(order_id)
        ON DELETE CASCADE ON UPDATE CASCADE
);
```

Table Relationship:



Data Flow:

1. Extract: Read 3 tables from RDS MySQL via JDBC connector
2. Transform: Perform inner joins on customer_id and order_id to create denormalized view
3. Load: Write to Keyspaces using Spark Cassandra Connector

Key Configuration:

- SSL enabled for Keyspaces connection (port 9142)
- IAM service-specific credentials for authentication
- Truststore configured for SSL certificate validation

Output Schema:

- Partition Key: customer_id (int)
- Clustering Key: order_id (int)
- Additional columns: name, email, city, order_date, amount, item_id, product_name, quantity

The screenshot shows the AWS Amazon Keyspaces CQL editor interface. On the left, there's a sidebar with links like 'Dashboard', 'Keyspaces', 'Tables', 'CQL editor', and 'Configuration'. The main area has tabs for 'Table view' and 'JSON view', with 'Table view' selected. A status bar at the top says 'Complete Execution time: 39 ms'. Below that, it says 'Records returned (50)'. A search bar says 'Find resources'. At the bottom right, there's a button 'Download results to CSV'. The table itself has columns: customer_id, order_id, amount, city, email, item_id, name, order_date, product_name, and quantity. The data consists of 50 rows of purchase details.

2. Read Keyspaces → Write Parquet to S3

Purpose: Extract denormalized sales data from Amazon Keyspaces and export to S3 as partitioned Parquet files for analytics and data lake storage.

Architecture:

- **Source:** Amazon Keyspaces table `sales_data` in keyspace `spark_keyspace`
- **Processing:** Apache Spark (Scala) with Cassandra Connector
- **Target:** AWS S3 bucket as partitioned Parquet files (`s3://sparkdemo-bucket-1/sales/parquet/`)

Data Flow:

1. **Extract:** Read `sales_data` from Keyspaces using Spark Cassandra Connector
2. **Transform:** Select subset of columns (`customer_id`, `order_id`, `amount`, `product_name`, `quantity`)
3. **Load:** Write to S3 as Parquet format, partitioned by `customer_id`

Key Configuration:

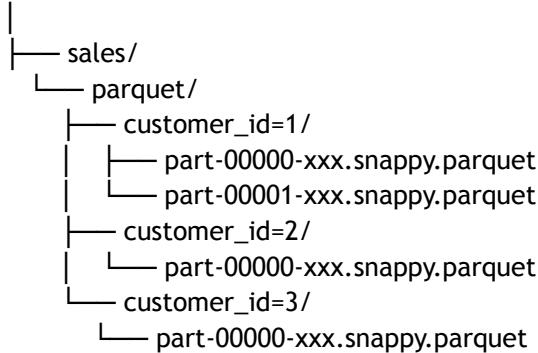
- SSL enabled for Keyspaces connection (port 9142)
- S3A file system with SimpleAWSCredentialsProvider for S3 access
- Extended timeouts for Keyspaces latency handling

Output Structure:

- Format: Parquet (columnar, compressed)
- Partitioning: By customer_id for optimized query performance
- Location: s3a://sparkdemo-bucket-1/sales/parquet/customer_id=X/

Directory Structure:

s3://sparkdemo-bucket-1/



The screenshot shows the AWS S3 console interface. The left sidebar navigation bar includes 'Amazon S3', 'Buckets', 'General purpose buckets', 'Table buckets', 'Vector buckets', 'Access management and security', 'Access Points', 'Access Points for FSx', 'Access Grants', 'IAM Access Analyzer', 'Storage management and insights', 'Storage Lens', 'Batch Operations', and 'Account and organization settings'. The main content area displays the 'parquet/' folder within the 'sales/' bucket. The 'Objects' tab is selected, showing 11 objects. The table lists the objects with columns for Name, Type, Last modified, Size, and Storage class. The objects are:

Name	Type	Last modified	Size	Storage class
customer_id=1/_SUCCESS	Folder	December 4, 2025, 01:29:17 (UTC+05:30)	0 B	Standard
customer_id=10/_	Folder	-	-	-
customer_id=2/_	Folder	-	-	-
customer_id=3/_	Folder	-	-	-
customer_id=4/_	Folder	-	-	-
customer_id=5/_	Folder	-	-	-
customer_id=6/_	Folder	-	-	-
customer_id=7/_	Folder	-	-	-
customer_id=8/_	Folder	-	-	-
customer_id=9/_	Folder	-	-	-

3. Read Parquet → Aggregate → Write JSON

Purpose: Read partitioned Parquet sales data from S3, compute product-level aggregations, and export results as JSON for downstream reporting and API consumption.

Architecture:

- **Source:** S3 Parquet files from Pipeline 2
(s3://sparkdemo-bucket-1/sales/parquet/)
- **Processing:** Apache Spark (Scala) with DataFrame aggregations
- **Target:** S3 JSON output (s3://sparkdemo-bucket-1/aggregates/products.json)

Data Flow:

1. **Extract:** Read partitioned Parquet files from S3 (customer_id partitions)
2. **Transform:** Aggregate sales data by product_name
 - Sum quantity → total_quantity
 - Sum amount → total_revenue
3. **Load:** Write aggregated results as single JSON file to S3

Aggregation Logic:

- **Group By:** product_name
- **Metrics:**
 - total_quantity = SUM(quantity)
 - total_revenue = SUM(amount)
- **Sort:** Descending by total_revenue (top revenue products first)

Output Format:

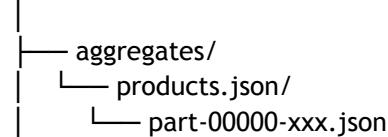
```
json
>{"product_name": "Widget A", "total_quantity": 150, "total_revenue": 45000.0}
>{"product_name": "Widget B", "total_quantity": 120, "total_revenue": 36000.0}
```

Key Configuration:

- S3A file system with SimpleAWSCredentialsProvider
- Configuration values stored as variables at top of code for easy management
- .coalesce(1) to produce single output file

Directory Structure:

s3://sparkdemo-bucket-1/



The screenshot shows the AWS S3 console interface. On the left, there is a navigation sidebar with the following structure:

- Amazon S3
- Buckets
 - General purpose buckets
 - Directory buckets
 - Table buckets
 - Vector buckets [New](#)
- Access management and security
 - Access Points
 - Access Points for FSx
 - Access Grants
 - IAM Access Analyzer
- Storage management and insights
 - Storage Lens
 - Batch Operations
- Account and organization settings
- AWS Marketplace for S3

The main panel displays the contents of the 'products.json/' folder. The 'Objects' tab is selected, showing two items:

Name	Type	Last modified	Size	Storage class
_SUCCESS	-	December 4, 2025, 02:10:05 (UTC+05:30)	0 B	Standard
part-00000-95be565c-bf9e-44a5-990b-99a025eb9992-c000.json	json	December 4, 2025, 02:09:58 (UTC+05:30)	3.8 KB	Standard

4. Spark Structured Streaming → Check RDBMS for New Records → Write to Kafka (Avro)

Purpose: Stream newly inserted order records from MySQL into Kafka as Avro messages for real-time processing and event-driven workflows.

Architecture:

- **Source:** MySQL table new_orders
- **Processing:** Spark Structured Streaming (Scala), JDBC incremental scan
- **Target:** Kafka topic orders_avro_topic (Avro-encoded)

Data Flow:

- **Extract:** Poll MySQL every 5 seconds using incrementColumn = order_id

- **Transform:** Convert each new row into Avro (using orders.avsc)
- **Load:** Publish Avro bytes to Kafka with order_id as the message key

Key Configuration:

- MySQL JDBC with incremental reading
- Spark trigger interval: 5 seconds
- Avro GenericRecord serialization
- Kafka producer using ByteArraySerializer

Avro Structure: orders.avsc

```
{
  "type": "record",
  "name": "OrderRecord",
  "namespace": "com.retail",
  "fields": [
    { "name": "order_id", "type": "int" },
    { "name": "customer_id", "type": "int" },
    { "name": "amount", "type": "double" },
    { "name": "created_at", "type": "string" }
  ]
}
```

Output Structure:

- **Format:** Avro binary
- **Topic:** orders_avro_topic
- **Key:** order_id
- **Value:** Avro payload containing order_id, customer_id, amount, created_at

```
Last login: Thu Dec  4 02:07:18 on ttys007
[rac0192 ~ % kafka-topics.sh --create --topic orders_avro_topic --bootstrap-server localhost:9092 --partitions 3 --replication-factor 1
WARNING: Due to limitations in metric names, topics with a period ('.') or underscore('_') could collide. To avoid issues it is best to use either, but not both.
Created topic orders_avro_topic.
[rac0192 ~ % kafka-topics.sh --list --bootstrap-server localhost:9092
_consumer_offsets
allocation.events
maintenance.events
notifications
orders_avro_topic
overdue.events
people-topic
reservation.notifications
[rac0192 ~ % kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic orders_avro_topic --from-beginning
2025-12-04 02:53:13
???\V@2025-12-04 02:53:13
02@2025-12-04 02:53:13
?b@2025-12-04 02:57:00
???(\V@2025-12-04 02:57:00
02@2025-12-04 02:57:00
[
```

5. Kafka Consumer Stream → Write JSON to S3

Purpose: Consume Avro messages from Kafka in real-time, decode them, and stream as JSON to S3.

Architecture:

- **Source:** Kafka topic orders_avro_topic (Avro binary from Pipeline 4)
- **Processing:** Spark Structured Streaming with Avro deserialization
- **Target:** S3 JSON files (s3://sparkdemo-bucket-1/stream/json/)

Data Flow:

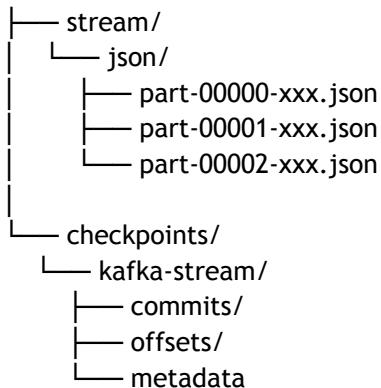
1. **Consume:** Read from Kafka topic in real-time
2. **Decode:** Deserialize Avro binary using schema (order_id, customer_id, amount, created_at)
3. **Load:** Write JSON to S3 every 10 seconds

Key Configuration:

- Starting offsets: earliest (reads all messages)
- Trigger: 10-second micro-batches
- Checkpoint: /tmp/kafka-s3-checkpoint (fault tolerance)
- Output mode: Append

Directory Structure:

s3://sparkdemo-bucket-1/



Output Format:

The screenshot displays two views of the same dataset. The top part is a terminal window showing JSON output from a command like `racit > Downloads > part-00000-029cd60b-b66d-4d89-af9f-4f829aa71966-c000.json`:

```
{ "order_id":13,"customer_id":1,"amount":158.75,"created_at":"2025-12-04 11:03:48","kafka_timestamp":"2025-12-04T11:07:09.915+05:30"}  
{ "order_id":14,"customer_id":2,"amount":89.99,"created_at":"2025-12-04 11:03:48","kafka_timestamp":"2025-12-04T11:07:09.922+05:30"}  
{ "order_id":15,"customer_id":3,"amount":428.0,"created_at":"2025-12-04 11:03:48","kafka_timestamp":"2025-12-04T11:07:09.922+05:30"}
```

The bottom part is the AWS S3 console under the bucket `sparkdemo-bucket-1` and folder `stream/json/`. It shows four objects:

Name	Type	Last modified	Size	Storage class
_spark_metadata/_	Folder	-	-	-
part-00000-029cd60b-b66d-4d89-af9f-4f829aa71966-c000.json	json	December 4, 2025, 11:07:20 (UTC+05:30)	397.0 B	Standard
part-00000-7d6f3009-b500-48c3-bf15-c72db67249a1-c000.json	json	December 4, 2025, 11:08:04 (UTC+05:30)	397.0 B	Standard
part-00000-eae4c6eb-a9f5-400d-b3f9-19877dd2050d-c000.json	json	December 4, 2025, 11:07:09 (UTC+05:30)	397.0 B	Standard